

# LUCID++: Enhanced Deep-Learning Solution for DDoS Attack Detection

Syed M Minhaj Naqvi

Department of Computing and Mathematics, KFUPM, Dhahran, Saudi Arabia.

Corresponding author(s). E-mail(s): [g202426160@kfupm.edu.sa](mailto:g202426160@kfupm.edu.sa);

## Abstract

Distributed Denial of Service (DDoS) attacks are some of the most pervasive and consequential threats to the availability of online services. Although the lightweight deep learning model architectures, like LUCID, offer promising avenue for fast and low resource-hungry detections, their applicability on newer, evolved, and more complex datasets is largely unexplored. This paper addresses this gap by doing evaluation of the original LUCID architecture on the CIC-DDoS2019 dataset, which is a modern benchmark containing sophisticated reflective and amplification attacks that were not incorporated in previous studies. Additionally, we put forth an enhanced model, LUCID++, which includes Batch Normalization and a hidden dense layer in order to improve the training and classification performance. Both the baseline as well as the enhanced models were thoroughly tuned via a 5-fold cross-validation process using `GridSearchCV`. The baseline model was able to achieve a strong accuracy of 98.50%, but our enhanced version surpassed this benchmark, reaching 98.93% accuracy and an F1-score of 0.9903. Crucially, this performance gain comes from a more optimized decision boundary which reduces the False Positive Rate by almost 50% while keeping a 99.22% True Positive Rate. These outcomes reinforce that the lightweight 1D-CNN design still remains a largely effective means for detecting modern DDoS attacks. The results also confirm that targeted architectural refinements can give significant practical improvements in the reliability of detection.

**Keywords:** DDoS Detection, Deep Learning, CNN, LUCID, Cybersecurity

## 1 Introduction

Distributed Denial of Service (DDoS) attacks continue to be one of the most damaging and critical threats to the network availability. Such attacks have the capability to overwhelm and paralyze critical online services along with the potential to cause catastrophic loss of availability to vital systems, including Critical National Infrastructure (CNI). It is a big challenge to detect these attacks quickly as detection systems must analyze a gargantuan volume of network traffic, all in real-time. This issue is only exacerbated by the continual evolution of attack vectors, rendering static and rule-based defense techniques to be very ineffective. Thus, there is an urgent need for detection techniques which are not only extremely accurate but also very computationally efficient for operating at line speed.

To address this imperative, lightweight Deep Learning (DL) solutions have shown significant promise. A notable example is LUCID, a DL solution that utilizes a streamlined Convolutional Neural Network (CNN) to achieve state-of-the-art detection accuracy with a processing overhead that is 40 times lower than competing heavy models [1]. The original LUCID study validated this technique across several datasets, namely ISCX2012, CICIDS2017, and CSE-CIC-IDS2018, demonstrating its viability in resource-constrained environments.

Nonetheless, the threat landscape is always evolving and existing models' high performance on older data is not a guarantee of success on newer, more complex, and evolved attack patterns. The CIC-DDoS2019 dataset by UNB, featuring huge traffic and varying set of modern reflective and amplification-based attacks, was not part of the original LUCID study. Hence, its performance on this contemporary CIC-DDoS2019 dataset remains an open question.

This paper aims to bridge this gap through a threefold contribution. Firstly, we establish a comprehensive performance baseline by implementing and evaluating the original LUCID architecture on the CIC-DDoS2019 dataset. Secondly, we propose and implement an enhanced version of the original model architecture, named LUCID++, which integrates Batch Normalization and an additional dense layer to augment classification capability and training stability. Lastly, we present a comprehensive performance evaluation of this enhanced model against the baseline. Our results demonstrate that while the original model achieves a strong accuracy of 98.5%, our enhanced model improves this to 98.93%. Most notably, it achieved a reduction of nearly 50% in the False Positive Rate on the test set, a critical metric for reducing alert fatigue in operational environments.

The rest of this paper is: Section 2 details the related work in the field. Section 3 details the methodology, including the dataset-specific characteristics, preprocessing steps, and model architectures. Section 4 details the experimental results and analysis, Section 5 provides a discussion of the findings, and Section 6 concludes with future work based on our results.

## 2 Background and Related Work

The detection of Distributed Denial of Service (DDoS) attacks is a well-established field within cybersecurity research. Over time, methodologies have shifted from simple statistical tests to high-performance deep learning models capable of automated feature extraction.

### 2.1 Traditional detection paradigms

Early work on this problem employed statistical as well as threshold-centric detection techniques [2, 3]. These early systems were engineered to monitor the network traffic for irregular patterns in packet headers, such as changes in the entropy of source IP addresses [4]. The logic behind this was, a volumetric DDoS attack would cause a sudden deviation from normal traffic behavior. This approach was not really robust as selection of effective threshold across different network environments was quite difficult [5]. Attackers could still bypass them via different techniques like using low-and-slow or multi-vector strategies.

### 2.2 Classical machine learning in intrusion detection

Limits of fixed thresholds prompted a shift toward classical Machine Learning (ML). This paradigm moved the focus from static rules to data-driven classification. Buczak and Guven survey common models such as Support Vector Machines (SVMs), k-Nearest Neighbors (k-NN), and Random Forests for intrusion detection [6]. While these models successfully raised detection accuracy, their performance remained very dependent on extensive, meticulous, and time-consuming feature engineering. Domain experts were needed for manually designing features like flow duration and packet inter-arrival times, consequently creating a bottleneck that slowed adaptation to novel attack types.

### 2.3 Deep learning for DDoS detection

In recent times, Deep Learning (DL) has come out as the dominant technique because it learns features directly from raw data with minimal to no manual intervention [1]. These models ingest low-level packet inputs and automatically learn very complex and non-linear patterns that distinguish benign traffic from malicious.

Two architectures have become dominant in this domain:

- Recurrent Neural Networks (RNNs), particularly Long Short-Term Memory (LSTMs) are used to model the temporal sequence of packets in a flow, effectively capturing time-based attack patterns [7].
- Convolutional Neural Networks (CNNs) treat traffic flows as spatial structures, like 1D or 2D images, and learn from the spatial relationships between packet features within a flow [8].

## 2.4 The lightweight model imperative and LUCID

While effective, many deep learning models are computationally heavy, containing millions of trainable parameters. DeepDefense, which combined CNNs and RNNs, is one such example [7]. This high level of complexity raises processing cost and hinders real-time deployment on resource-constrained edge gateways or on-device IoT platforms [9, 10].

This gap in performance motivates the need for lightweight and accurate models. LUCID (Lightweight, Usable CNN in DDoS Detection) is a primary example of this philosophy [1]. Doriguzzi-Corin et al. designed LUCID as a practical deep model with exceptionally low overhead. Their network is a compact 1D-CNN with only 2,241 parameters. The authors showed that it matches the performance of much heavier models like 3LSTM, which has over one million parameters, while running approximately 40 times faster.

LUCID's combination of accuracy and speed makes it a good baseline for modern DDoS detection. We use LUCID as the baseline for our work as well, validating it against the CIC-DDoS2019 dataset which the original study did not include, and trying architectural enhancements for pushing the detection performance even beyond the baseline model.

## 3 Methodology

This study employs a comparative experimental design to rigorously evaluate the efficacy of lightweight deep learning models. We test the original LUCID model against our improved version, LUCID++. Both models are trained and evaluated on the CIC-DDoS2019 dataset, a modern benchmark not included in the original LUCID study. The primary objective is to quantify how specific architectural changes, namely the addition of Batch Normalization and a dense hidden layer, influence training stability, convergence speed, and overall detection accuracy.

### 3.1 Dataset and Preprocessing

The CIC-DDoS2019 dataset was selected for its relevance to the current threat landscape. It contains a broad mix of attack types, including sophisticated reflective and amplification-based flood attacks such as NTP, DNS, LDAP, and NetBIOS. It provides modern traffic patterns that reflect current DDoS techniques far better than older datasets.

The preprocessing pipeline strictly follows the lightweight methodology described in the original LUCID paper to ensure fair comparison. Traffic is divided into non-overlapping ten-second windows, and packets within each window are grouped by bi-directional flow. Each flow is then split into samples containing up to ten packets. For every packet, eleven distinct features are extracted: relative timestamp, packet length, IP flags, protocol information, and transport-layer fields for TCP, UDP, and ICMP. Crucially, IP addresses and port numbers are excluded to ensure the model learns generalizable attack patterns rather than memorizing specific host identities. Flows shorter than ten packets are padded with zeros, and all features are normalized

to a scale between zero and one. The final data shape is a three-dimensional tensor of size  $(10, 11, 1)$ , representing ten packets, eleven features, and one channel, a format well-suited for two-dimensional convolution.


### 3.2 Model Architectures

Two neural models were built and trained using identical optimization and loss settings to isolate the impact of architectural changes. Both models utilize the AdamW optimizer and binary cross-entropy as the loss function.

#### 3.2.1 Baseline Model (LUCID)

The baseline model is a direct implementation of the original LUCID design [1]. As described in the paper, this architecture uses a convolutional layer where the filter's width,  $f$ , matches the width of the input feature matrix ( $f = 11$ ). This design forces the filter to slide in only one dimension (along the  $n$  packets), thus behaving as a 1D convolution over the temporal axis.

In our implementation, this is built using a 2D Convolutional Layer ('Conv2D') with a kernel size of  $(h, f)$ . This is followed by a ReLU activation, a 'GlobalMaxPooling2D' layer, a 'Flatten' operation, and a 'Dense' output layer with a sigmoid activation. The original architecture is depicted in Figure 1.




**Fig. 1** The original LUCID model architecture (adapted from [1]).

#### 3.2.2 Proposed Enhanced Model (LUCID++)

The proposed model, LUCID++, extends this design with two modifications. A **BatchNormalization** layer follows the 'Conv2D' layer to stabilize activations and improve convergence. A **Dense** hidden layer with  $n$  units and ReLU activation is

added after the ‘Flatten’ operation to learn richer combinations of extracted features. Dropout regularization is applied to this new layer to prevent overfitting. This enhanced architecture is shown in Figure 2.



**Fig. 2** The proposed LUCID++ model architecture.

### 3.3 Evaluation Setup

Model evaluation relies on a rigorous five-fold cross-validation process managed through ‘GridSearchCV’. Each fold trains and validates models across a pre-defined hyperparameter grid, and the configuration yielding the highest F1-score on the validation folds is selected for retraining on the full training set. The grid explores a range of values for learning rate (0.01 and 0.001), batch size (1024 and 2048), convolution filter count (32 and 64), L1 regularization, convolutional dropout (0 or 0.2), and AdamW weight decay (0.01, 0.001, and 0). For the LUCID++ model, the size of the dense layer and its dropout rate are also varied.

The final, optimized model is tested on a strictly held-out subset of the dataset to ensure unbiased evaluation. Performance is measured using a suite of metrics including accuracy, F1-score, true positive rate, false positive rate, false negative rate, true negative rate, and the ROC–AUC curve. These metrics provide a holistic view of the model’s detection reliability and error balance.

### 3.4 Implementation and Reproducibility

All experiments were done on a Windows 11 workstation having an NVIDIA 4060 GPU and 32GB of RAM. The models were implemented in Python 3 using TensorFlow 2.x and the Keras API. The Scikeras library was used to interface Keras models with Scikit-learn for hyperparameter search. Data handling and I/O was managed by NumPy and h5py. Additionally, PyShark was used for parsing raw PCAP data, requiring the installation of T-shark (comes with Wireshark).

For ensuring the reproducibility of our results, a fixed random seed of 42 was set across NumPy, TensorFlow, and Scikit-learn. This guarantees consistency in data splits

and model initialization for each run. The complete source code, the model definitions, and all the hyperparameter settings are publicly available online in the GitHub repository for this project [11], allowing for exact replication of our experiments.

## 4 Results and Analysis

This section gives a detailed account of our experimental results that were gathered during our study. We start by elaborating the extensive hyperparameter tuning process, then proceed to a deep dive into the training dynamics and convergence behavior of the final model, and end with a rigorous evaluation of the classification performance on the unseen test set.

### 4.1 Hyperparameter Tuning and Model Performance

The optimization of the model’s hyperparameter was done using a comprehensive 5-fold cross-validation technique, orchestrated via the ‘GridSearchCV’ framework. This robust validation procedure was applied to the whole entirety of the training dataset in order to ensure that the selected parameters were not biased towards any specific subset of the data. In total, the search space covered 1,920 unique combinations of hyperparameters. Each combination was separately trained and evaluated, ensuring that the final configuration selected was truly the global optimum in our defined search space.

The search process identified the following specific configuration as optimal for maximizing the F1-score of the enhanced LUCID++ model:

- **Batch Size:** 1024, A larger batch size was favored, likely providing more stable gradient estimates during updates.
- **Learning Rate:** 0.001, A more conservative learning rate was selected, allowing for finer adjustments to the weights as the model converged.
- **Weight Decay (AdamW):** 0.001, The inclusion of weight decay indicates that regularization played a key role in preventing overfitting.
- **Kernels (Filters):** 64, The higher number of filters suggests the model benefited from a richer feature extraction capacity.
- **Convolutional Dropout:** None, Interestingly, dropout was not required at the convolutional stage, suggesting the Batch Normalization provided sufficient regularization there.
- **Regularization (L1):** None, Explicit L1 penalty was not selected, deferring regularization duties to weight decay and dense dropout.
- **Dense Units:** 16, A relatively compact hidden layer was sufficient to capture the necessary non-linear relationships.
- **Dense Dropout:** 0.5, A high dropout rate in the dense layer confirms its role in forcing the model to learn redundant, robust representations.

Table 1 gives a comparative performance analysis between the original LUCID baseline model and our final, enhanced LUCID++ model on the held-out CIC-DDoS2019 test set. The data shows that the LUCID++ model achieved a quantifiable

improvement, getting an accuracy increase of 0.0043 (0.43%) and an F1-Score increase of 0.0037 (0.37%).


While these numerical improvements might appear marginal, their significance is profound when analyzing the balance between True Positives and False Positives. The baseline model was highly aggressive, achieving a near-perfect True Positive Rate (TPR) but at the expense of a False Positive Rate (FPR) of 2.86%. In contrast, our enhanced model achieved a better equilibrium. It successfully reduced the FPR by about 50%, bringing it down to 1.43%, while incurring only a 0.78% cost in the False Negative Rate (FNR). This shift shows a substantial practical improvement, as it drastically cuts down on the false alarms without compromising the system’s ability to detect actual threats.

**Table 1** Performance Comparison on the CIC-DDoS2019 Test Set

Model	Accuracy	F1-Score	TPR	FPR	FNR	TNR
Baseline (LUCID)	0.9850	0.9866	0.9961	0.0286	0.0039	0.9714
LUCID++	<b>0.9893</b>	<b>0.9903</b>	0.9922	0.0143	0.0078	0.9857

## 4.2 Training Dynamics

The training behavior of our LUCID++ architecture can be seen in Figure 3, which plots the accuracy and loss over the course of 300 training epochs. These curves provide critical insight into the stability as well as the learning progression of the model. Referring to the loss curves on the right side of Figure 3, we can observe a consistent and steady decrease in the loss value for both the training part (depicted in blue) and the validation part (depicted in orange) as the training epochs progress. Crucially, the validation loss curve closely follows the training loss curve throughout the entire process and does not show any kind of divergence or sustained increase. This behavior is a strong indicator that the regularization and optimization techniques we used, specifically the combination of Dense Dropout, Batch Normalization, and the AdamW optimizer, were highly effective in preventing the model from overfitting to the training data. The accuracy curves on the left show an interesting two-part learning process. Initially, there is a period of instability with high fluctuation, particularly evident in the validation accuracy, where the model explores the feature space and adjusts its weights from their random initializations. However, a sharp and a stable convergence is observed around epoch 100. Following this point, both the training and validation accuracy metrics climb rapidly to near-perfect levels and remain tightly coupled until the end of the training run. This close tracking confirms that the model has successfully learned a generalized representation of the complex data patterns, rather than simply memorizing specific examples from the training set.



**Fig. 3** Model Accuracy (Left) and Loss (Right) During Training (LUCID++ Architecture)

### 4.3 Classification Analysis

In order for providing a more granular and meticulous testing of the final model's performance, we generated a detailed classification report and a Receiver Operating Characteristic (ROC) curve using the unseen test dataset.


The classification report, shown in Table 2, breaks down the precision, recall, and F1-score for each specific class. The model showed exceptional performance, achieving scores of approximately 0.99 across all these three metrics for both the 'Benign' as well as the 'DDoS' class. Specifically, the model correctly identified 99.22% of all malicious DDoS attack samples (Recall/TPR), while simultaneously correctly identifying 98.57% of all benign traffic samples (Recall/TNR). This consistency in performance across classes confirms the balanced nature of the classifier, and confirms that it is not biased toward the majority class.

**Table 2** Classification Report for LUCID++ on Test Set

	Precision	Recall	F1-Score	Support
<b>Benign (0)</b>	0.9904	0.9857	0.9881	210
<b>DDoS (1)</b>	0.9884	0.9922	0.9903	258
<b>Accuracy</b>			<b>0.9893</b>	<b>468</b>
<b>Macro Avg</b>	0.9894	0.9890	0.9892	468
<b>Weighted Avg</b>	0.9893	0.9893	0.9893	468

Figure 4 shows the critical trade-off that occurs between the True Positive Rate and the False Positive Rate across all possible decision thresholds. Our final model achieved an Area Under the Curve (AUC) of **0.9927**, a value that is very close to the theoretical maximum of 1.0. This extremely high AUC score indicates an outstanding level of class separability, meaning that the model can distinguish between benign and malicious traffic with a very high confidence. Visually, the curve shows a sharp

and almost vertical rise toward the top-left corner of the plot. This confirms that the model is capable of achieving a very high True Positive Rate while maintaining an extremely low False Positive Rate, further supporting its suitability for deployment in high-security environments where minimizing false alarms is extremely important.



**Fig. 4** ROC Curve for the Enhanced Model (LUCID++) on the Test Set (AUC = 0.9927)

## 5 Discussion

The experimental results detailed in Section 4 provide a strong validation for the architectural enhancements put forth in our LUCID++ model. By achieving a final accuracy of 98.93% and an F1-Score of 0.9903, the enhanced model not only outperformed the baseline, but it also showed a significantly more robust and operationally desirable performance.

The most significant finding of this study lies in the analysis of the error types. The baseline LUCID model showed a high sensitivity identifying nearly all attacks with a 99.61% TPR. However, this sensitivity came at the cost of a 2.86% False Positive Rate. In the context of a vast-telemetry Security Operations Center (SOC), a 2.86% false alarm rate can translate to a lot of incorrect alerts, leading to "alert fatigue" in analysts. Our optimized LUCID++ model successfully addressed this critical issue. By reducing the FPR to 1.43%, a decrease of nearly 50% and while maintaining a 99.22% detection rate, the enhanced model offers a more practical solution for real-world deployment.

We attribute these performance gains to the two primary architectural modifications we performed. First, the addition of a **BatchNormalization** layer played a central role in stabilizing the training process. By normalizing the activations output by the convolutional layer, this mechanism smoothed the optimization process,

preventing the model from getting stuck in poor local minima during the early epochs which were very volatile as seen in Figure 3. Second, the addition of the hidden **Dense(16)** layer provided the network with increased representational capacity. Unlike the original architecture, which translated extracted features directly to the output, this intermediate layer allowed the model to learn more complex, non-linear combinations of the features before making a final classification decision.

It is worth noting that while our model achieved an F1-Score of 0.9903 on the CIC-DDoS2019 dataset, this value is not directly comparable to the 0.9987 F1-Score reported in the original LUCID paper for the CSE-CIC-IDS2018 dataset [1]. These datasets differ significantly in terms of attack vectors, traffic volume, and network topology. However, the fact that our lightweight model achieved such elite performance on a newer and a more complex dataset serves as a strong validation that the 1D-CNN approach is highly effective and deployable to modern threat landscapes.

Despite these successes, the study has its limitations. The primary limitation is the remaining 1.43% False Positive Rate. While a significant improvement, further reduction is needed for automated mitigation systems. Additionally, our evaluation was restricted to only a single new dataset, leaving the model’s performance on wider variety of network environments an open question. Finally, while the model architecture remains relatively lightweight in terms of parameter count, a rigorous measurement of inference latency and memory usage on edge hardware was outside the scope of this study.

## 6 Conclusion and Future Work

### 6.1 Conclusion

This paper presented a meticulous evaluation and enhancement of lightweight LUCID deep learning model for DDoS attack detection. We started by making a strong performance baseline of 98.5% accuracy using the original architecture on the modern CIC-DDoS2019 dataset. Building off of this, we introduced LUCID++, an enhanced architecture including Batch Normalization, an additional dense hidden layer, and optimization via AdamW. Through an extensive hyperparameter search, our enhanced model achieved a superior accuracy of 98.93% and an F1-Score of 0.9903. This performance improvement was driven by a near-50% reduction in false positives achieving a far better balance between sensitivity and specificity. These findings confirm that lightweight 1D-CNN architectures remain a powerful and an efficient tool for network defense and demonstrate that their performance can be significantly improved through targeted modern architectural refinements.

### 6.2 Future Work

Based on the findings of this study, we propose three directions for future research endeavors:

- **False Positive Reduction:** Future work should focus specifically on bringing down the remaining 1.43% FPR. Potential techniques for this include the use of

advanced data augmentation to better represent the benign class or the integration of a secondary specialized anomaly filter designed to double-check borderline classifications.

- **Cross-Dataset Generalization:** For further validating the robustness of the LUCID++ architecture it should be evaluated across a broader spectrum of datasets, including the original ISCX2012, CICIDS2017, and CSE-CIC-IDS2018 datasets. This would help to determine if the performance gains we observed are universal or specific to the dataset we used.
- **Edge Deployment Benchmarking:** An important next step is to measure the real-world computational cost of the enhanced model. This would involve deploying LUCID++ on a resource-constrained edge device, such as the NVIDIA Jetson TX2 used in the original study [1] to rigorously compare its inference latency and memory footprint against the baseline model.

## Declarations

- **Funding:** Not applicable.
- **Conflict of interest:** The authors declare that they have no conflict of interest.
- **Data availability:** The dataset analyzed during the current study is available in the CIC-DDoS2019 repository.
- **Code availability:** The code is available at the GitHub repository referenced in the text.

## References

- [1] Doriguzzi-Corin, R., Millar, S., Scott-Hayward, S., Martínez-del-Rincón, J., Siracusano, D.: Lucid: A practical, lightweight deep learning solution for ddos attack detection. *IEEE Trans. Netw. Serv. Manag.* **17**(2), 876–889 (2020)
- [2] Bojović, P., Bašićević, I., Ocovaj, S., Popović, M.: A practical approach to detection of distributed denial-of-service attacks using a hybrid detection method. *Comput. Elect. Eng.* **73**, 84–96 (2019)
- [3] Kalkan, K., Altay, L., Gür, G., Alagöz, F.: Jess: Joint entropy-based ddos defense scheme in sdn. *IEEE J. Sel. Areas Commun.* **36**(10), 2358–2372 (2018)
- [4] Feinstein, L., Schnackenberg, D., Balupari, R., Kindred, D.: Statistical approaches to ddos attack detection and response. In: Proc. DARPA Inf. Survivability Conf. Expo., p. 303 (2003)
- [5] Shah, S.B.I., Anbar, M., Al-Ani, A., Al-Ani, A.K.: Hybridizing entropy based mechanism with adaptive threshold algorithm to detect ra flooding attack in ipv6 networks. In: Computational Science and Technology, pp. 315–323. Springer, Singapore (2019)
- [6] Buczak, A.L., Guven, E.: A survey of data mining and machine learning methods

- for cyber security intrusion detection. *IEEE Commun. Surveys Tuts.* **18**(2), 1153–1176 (2016)
- [7] Yuan, X., Li, C., Li, X.: Deepdefense: Identifying ddos attack via deep learning. In: Proc. SMARTCOMP, pp. 1–8 (2017)
  - [8] Wu, K., Chen, Z., Li, W.: A novel intrusion detection model for a massive network using convolutional neural networks. *IEEE Access* **6**, 50850–50859 (2018)
  - [9] Roopak, M., Tian, G.Y., Chambers, J.: Deep learning models for cyber security in iot networks. In: Proc. IEEE CCWC, pp. 452–457 (2019)
  - [10] Wang, N., Varghese, B., Matthaiou, M., Nikolopoulos, D.S.: Enorm: A framework for edge node resource management. *IEEE Trans. Services Comput.* (2017)
  - [11] Naqvi, S.M.M.: Github Repository for LUCID++: Enhanced Deep-Learning Solution for DDoS Attack Detection. <https://github.com/BRAIN-Lab-AI/LUCID-Enhanced-Deep-Learning-Solution-for-DDoS-Attack-Detection>